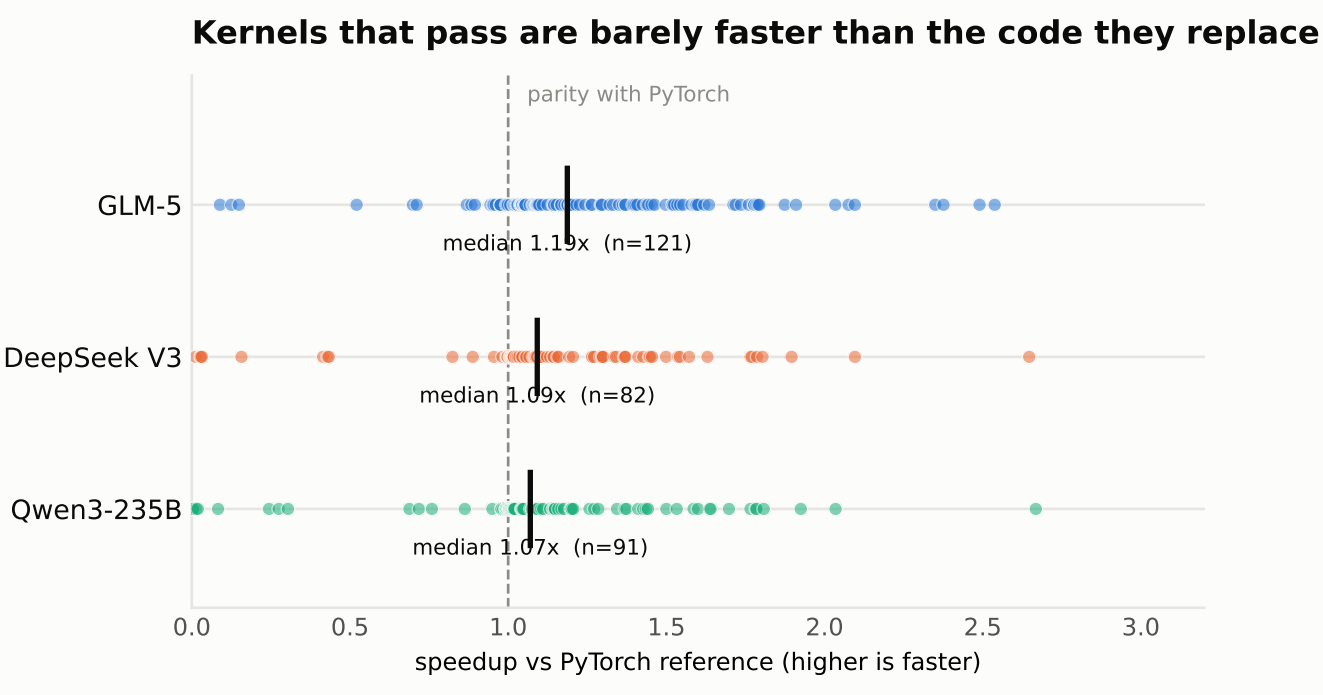
Kernels that pass are barely faster than the code they replace
parity with PyTorch
GLM-5
median 1.19x (n=121)
DeepSeek V3
median 1.09x (n=82)
Qwen3-235B
median 1.07x (n=91)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
speedup vs PyTorch reference (higher is faster)