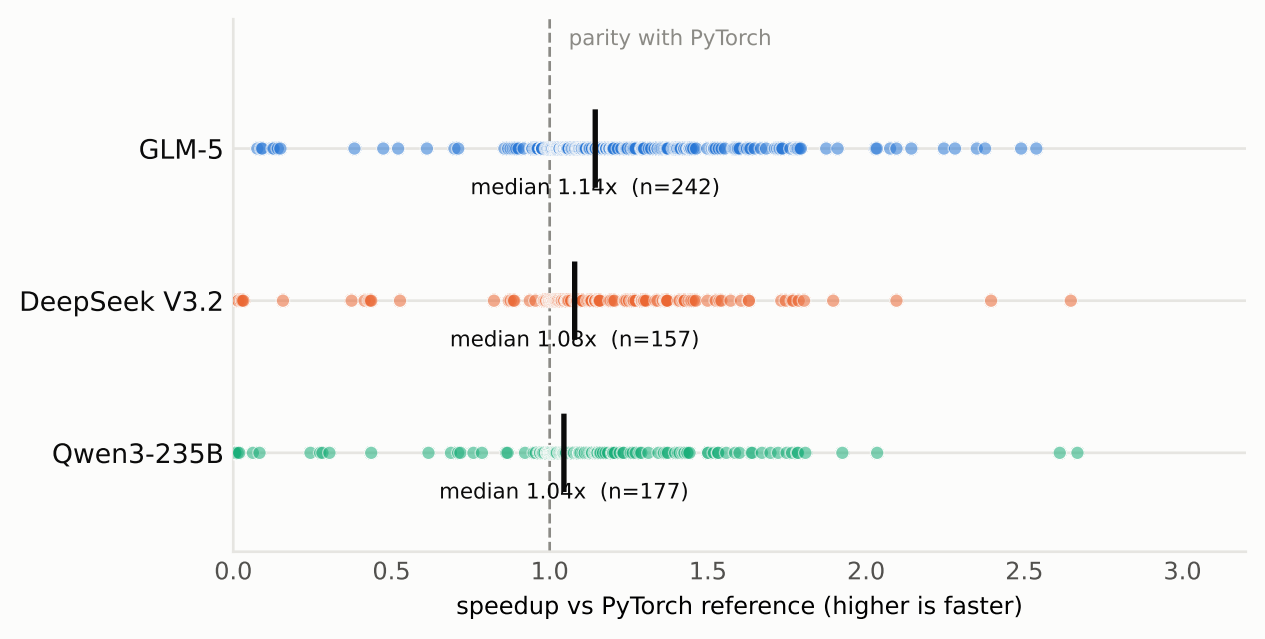
parity with PyTorch
GLM-5
median 1.14x (n=242)
DeepSeek V3.2
median 1.08x (n=157)
Qwen3-235B
median 1.04x (n=177)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
speedup vs PyTorch reference (higher is faster)